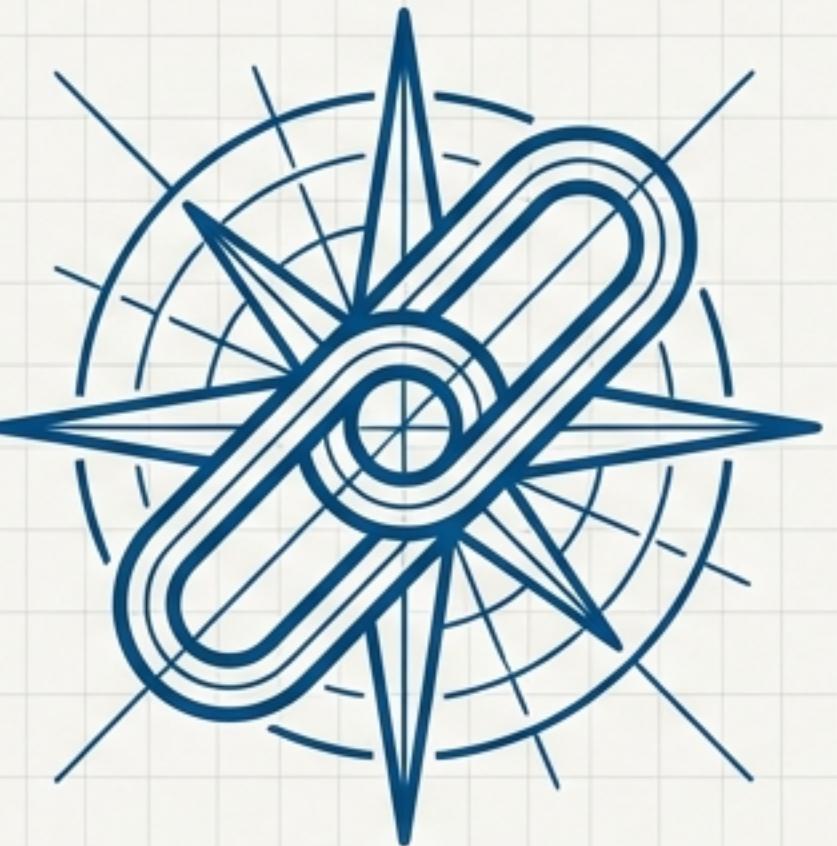


# The Architect's Journey: Designing a URL Shortener from the Ground Up



A step-by-step guide to building a scalable,  
high-performance service.

# Why do we turn long, unwieldy links...

https://www.example.com/products/category-x/item-12345?sessionid=abc-123-xyz-789&source=email&utm\_campaign=q4\_promo&utm\_content=banner\_ad

The URL is shown with callouts pointing to its parts: 'Domain' (the www.example.com part), 'Path' (the products/category-x/item-12345 part), 'Query Parameters' (the ?sessionid=abc-123-xyz-789&source=email&utm\_campaign=q4\_promo&utm\_content=banner\_ad part), and 'Session ID' (the sessionid value). Below the URL, arrows point from the query parameters to 'Source', 'Campaign', and 'Content' categories.

...into short,  
memorable ones?

**brand.ly/new-promo**

Custom Domain



Memorable Alias

Clean & Shareable

# A shorter link is a smarter link.



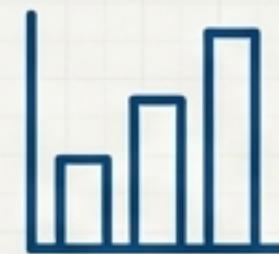
## More Memorable & Branded

Polished links like 'brand.ly/feature' are easier to remember and extend your brand's reach.



## Increased Click-Through

Clean, trustworthy links are more visually appealing and receive more clicks than long, complex URLs.



## Trackable Performance

Integrated analytics track click counts and sources, providing valuable engagement data.

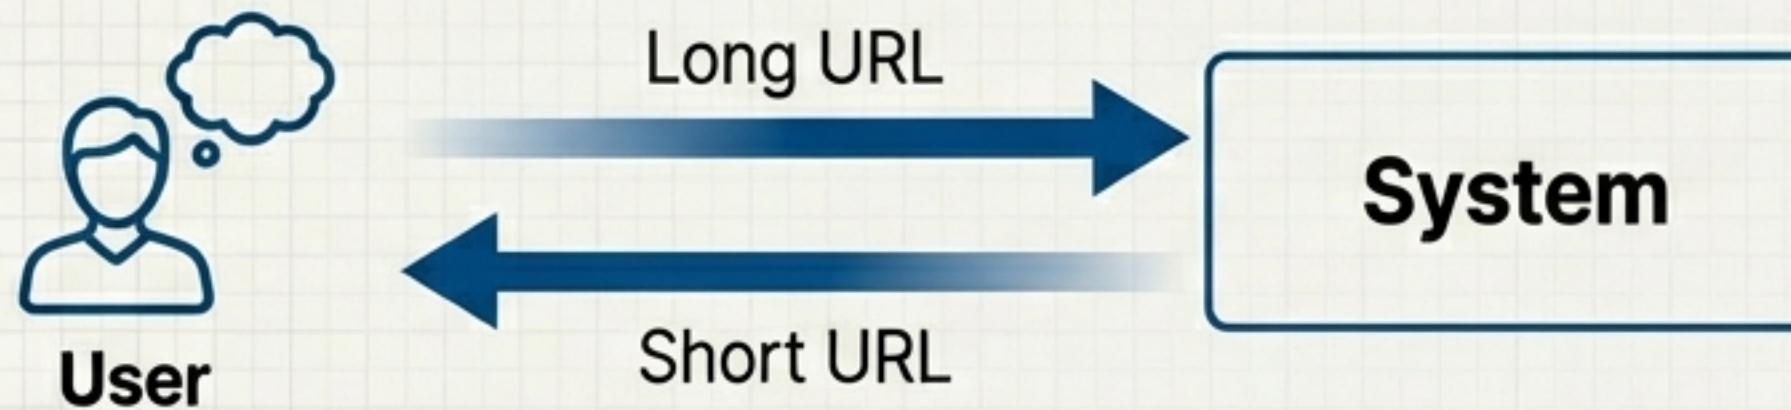


## Character-Efficient

Essential for character-limited platforms, freeing up space for your message.

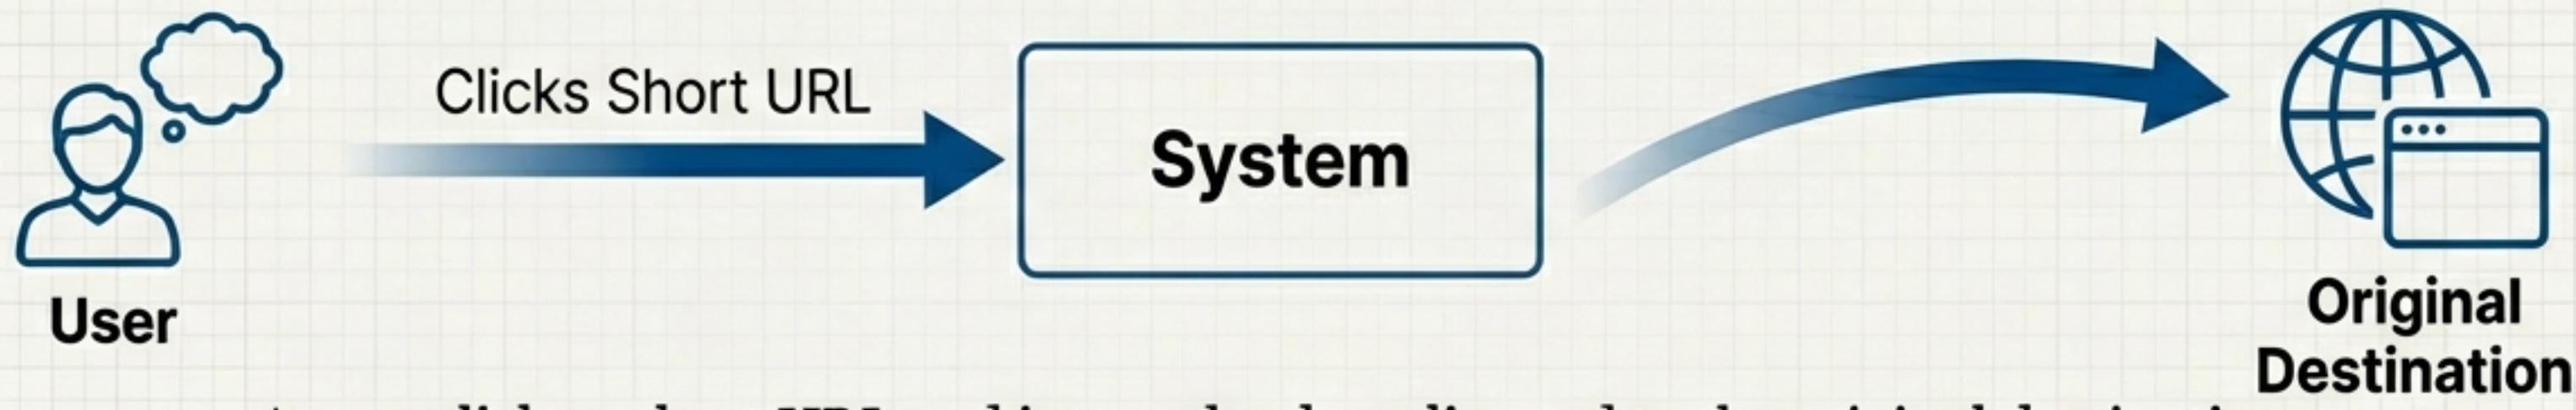
# What are the system's core functions?

## 1. Shorten:



A user provides an original URL and receives a shortened version.

## 2. Redirect:



A user clicks a short URL and is seamlessly redirected to the original destination.

# Beyond function: What defines success?



## High Availability:

99.9% uptime for redirects. Our service must always be on.



## High Throughput:

Capable of handling a massive volume of requests.



## Low Latency:

Redirects must be near-instantaneous to ensure a good user experience.

# How big is our user base?

**1,000,000**



We will start with an assumption of **1 Million Monthly Active Users (MAU)**.  
This single figure will drive all our estimations for traffic and storage.

# How many new links will we generate?



1M MAU



50% of users

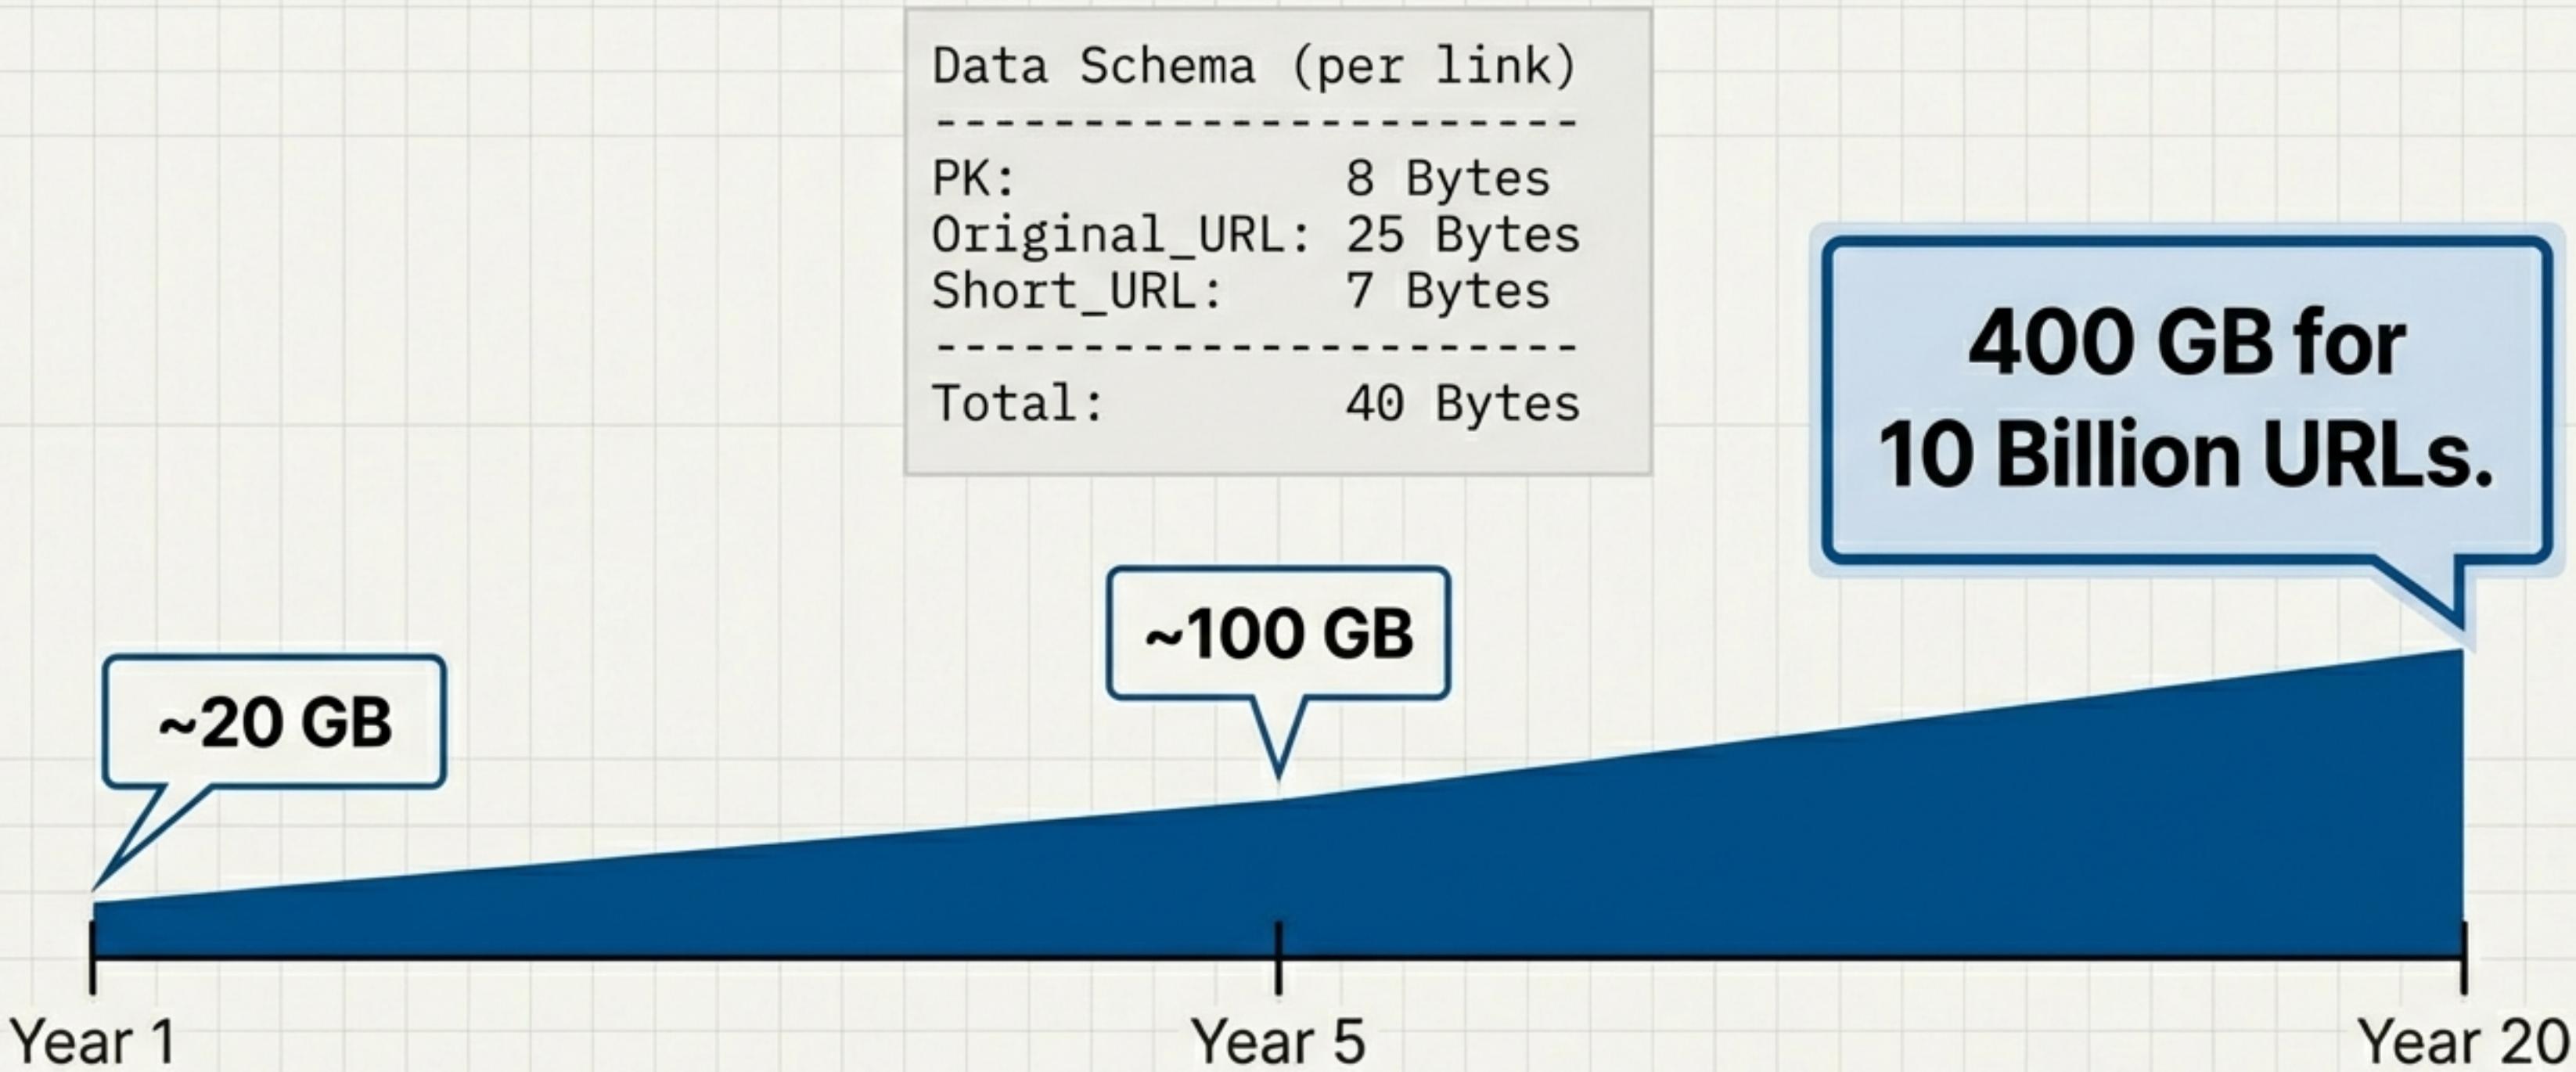


create 2 URLs/day

**1 Million new URLs created every single day.**

This translates to ~500 Million new URLs per year.

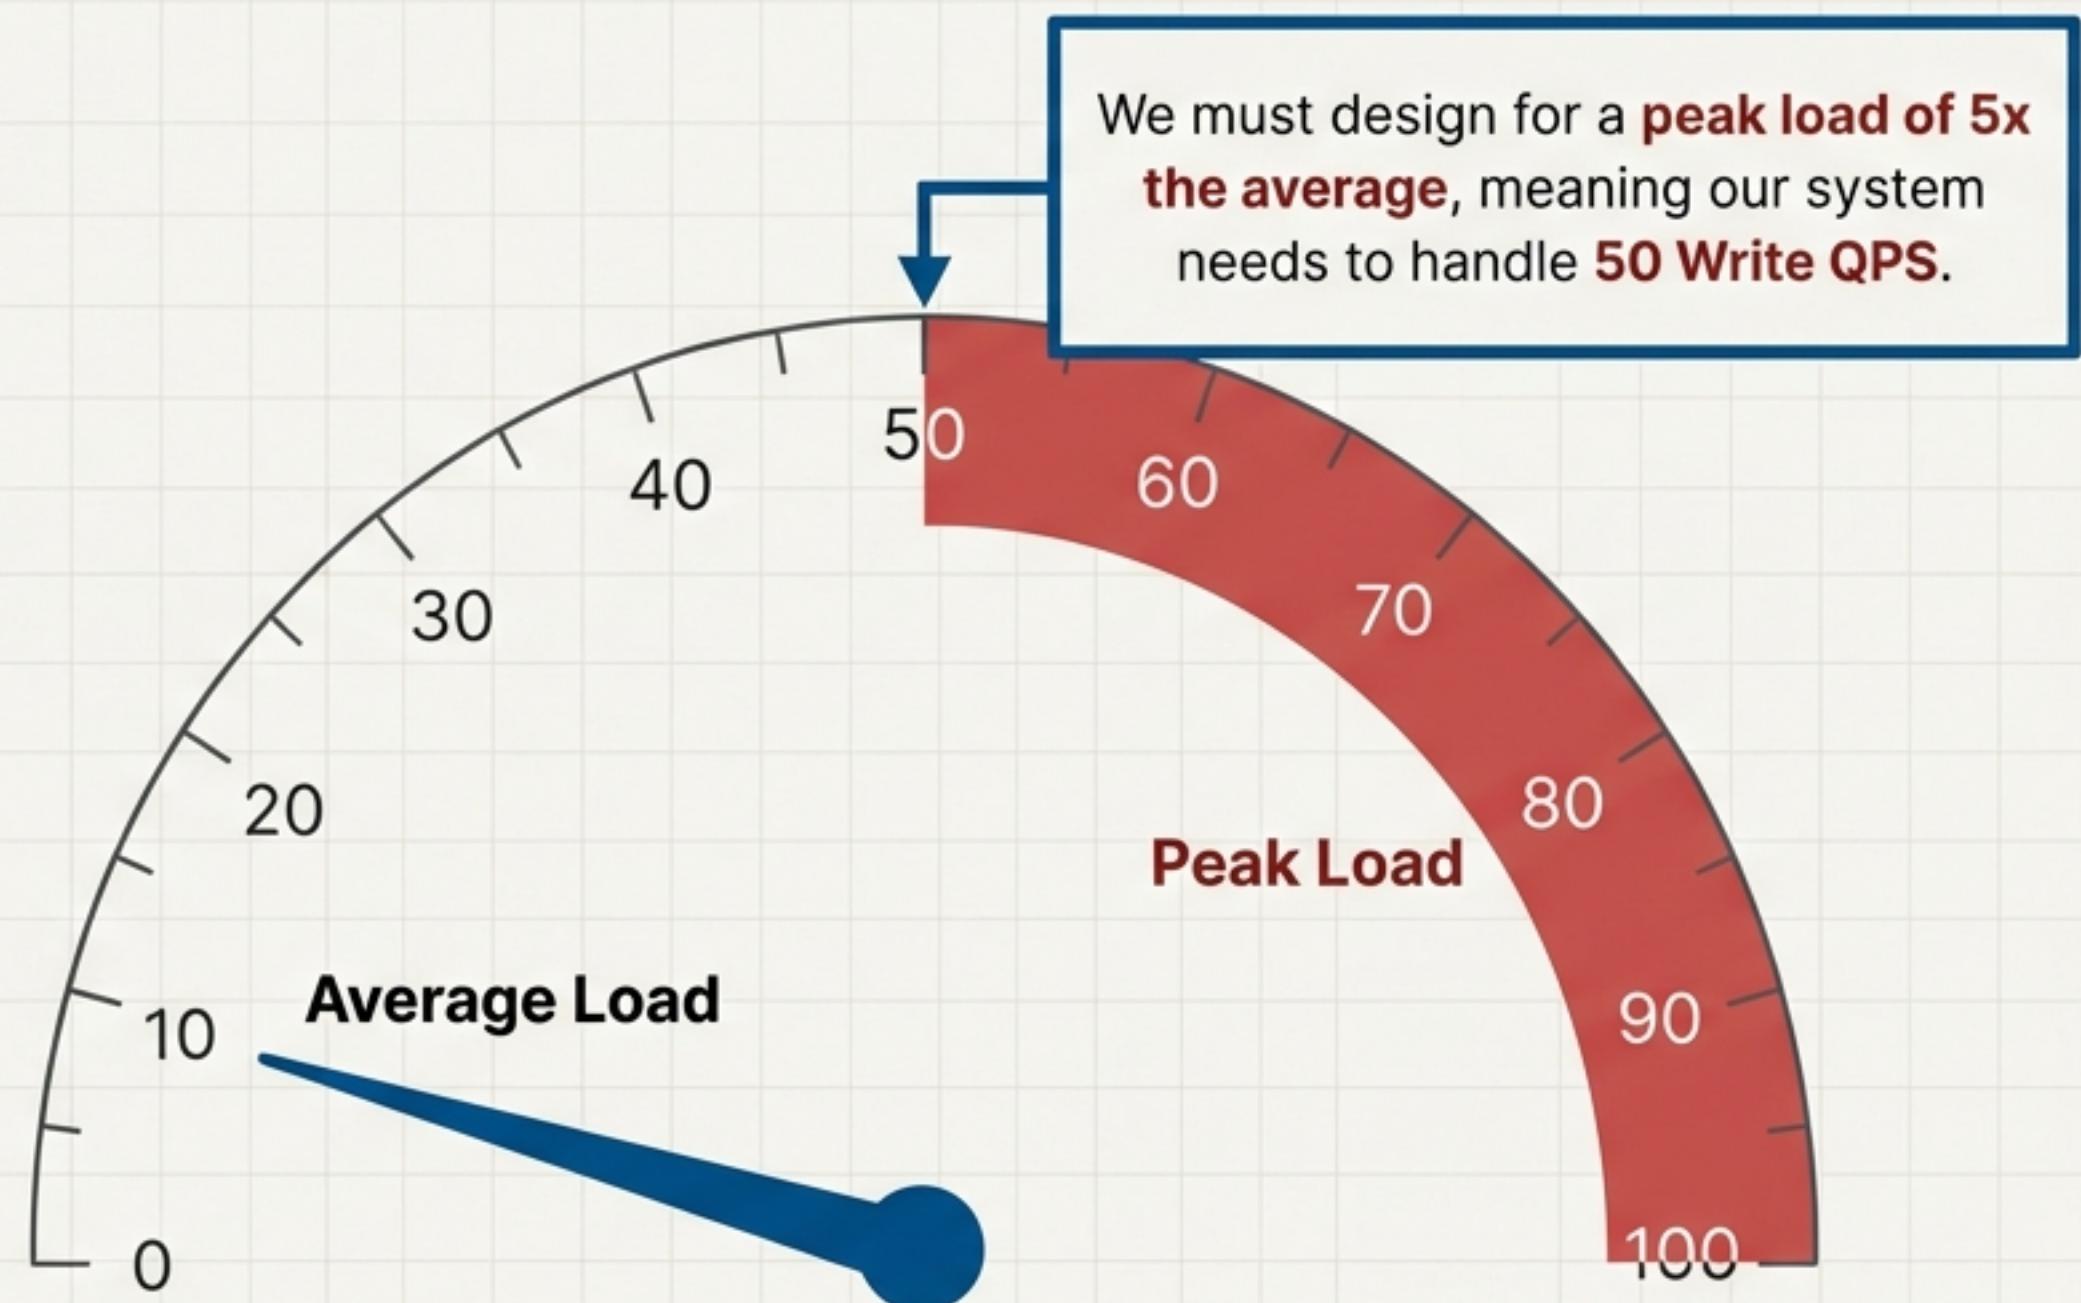
# How much storage will we need for 20 years?



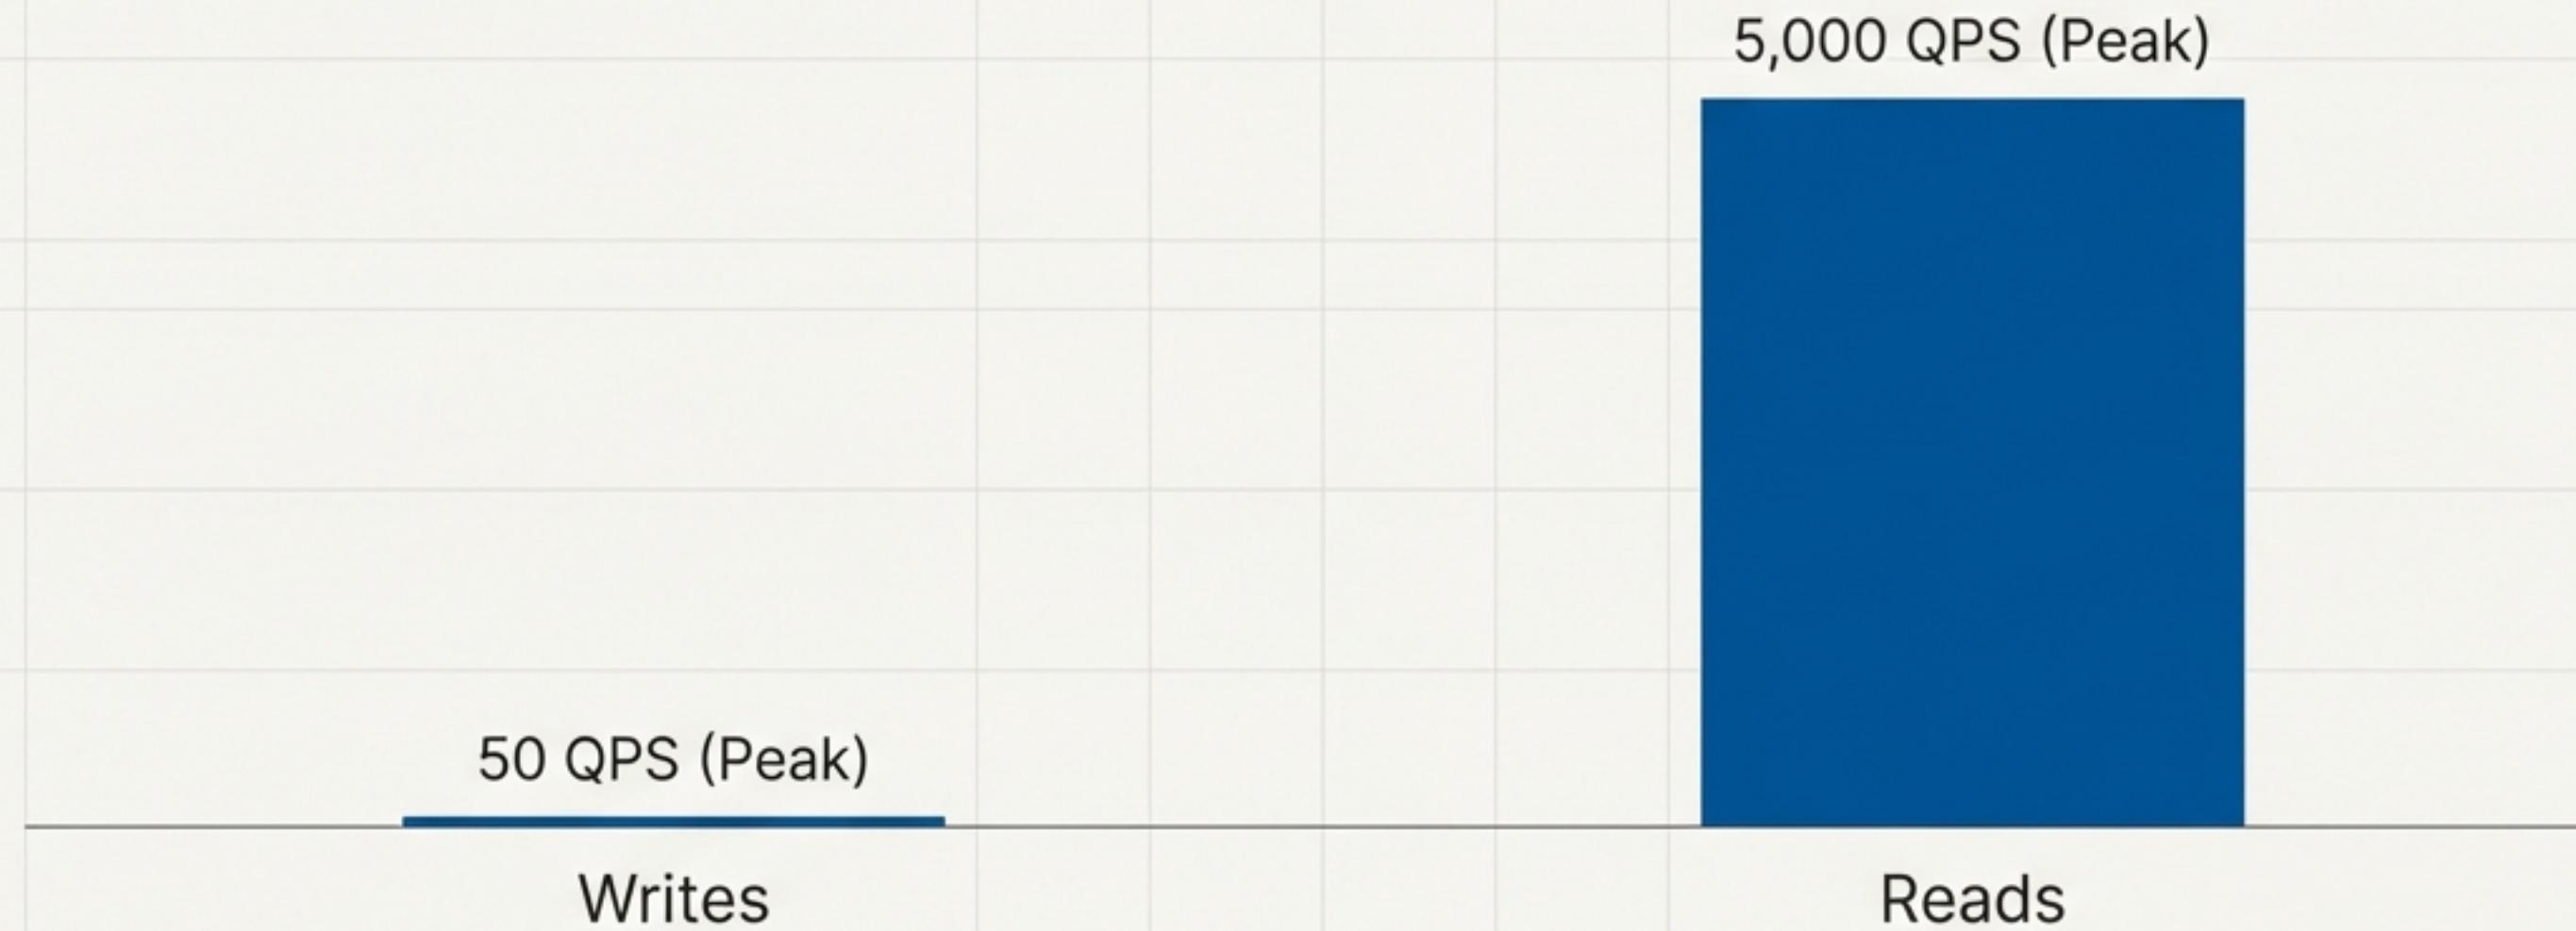
# How fast must our ‘write’ operations be?

1 Million URLs /  
~100,000 seconds per day

**~10 Write QPS  
(Average)**



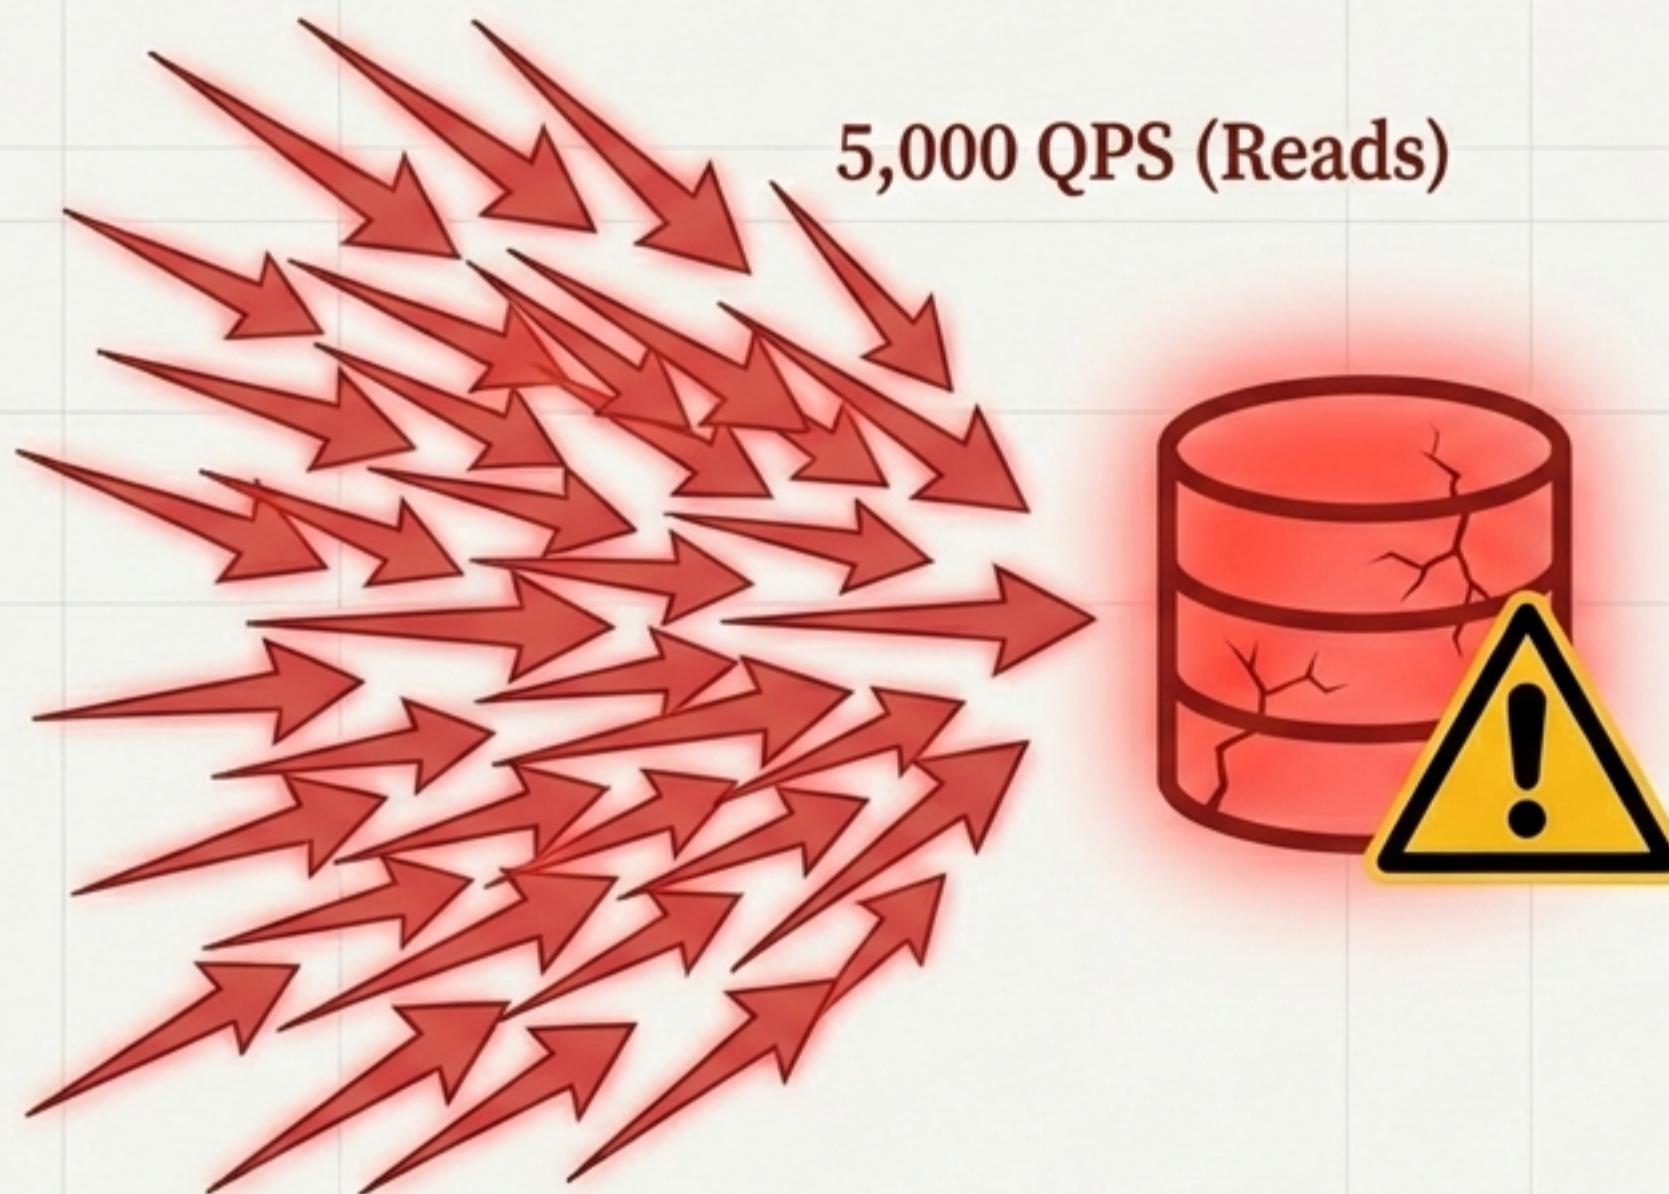
# The real challenge isn't writes. It's reads.



**A 100:1 Read-to-Write Ratio.**

For every link created, we anticipate 100 redirect requests.

# What happens if every redirect hits our database?

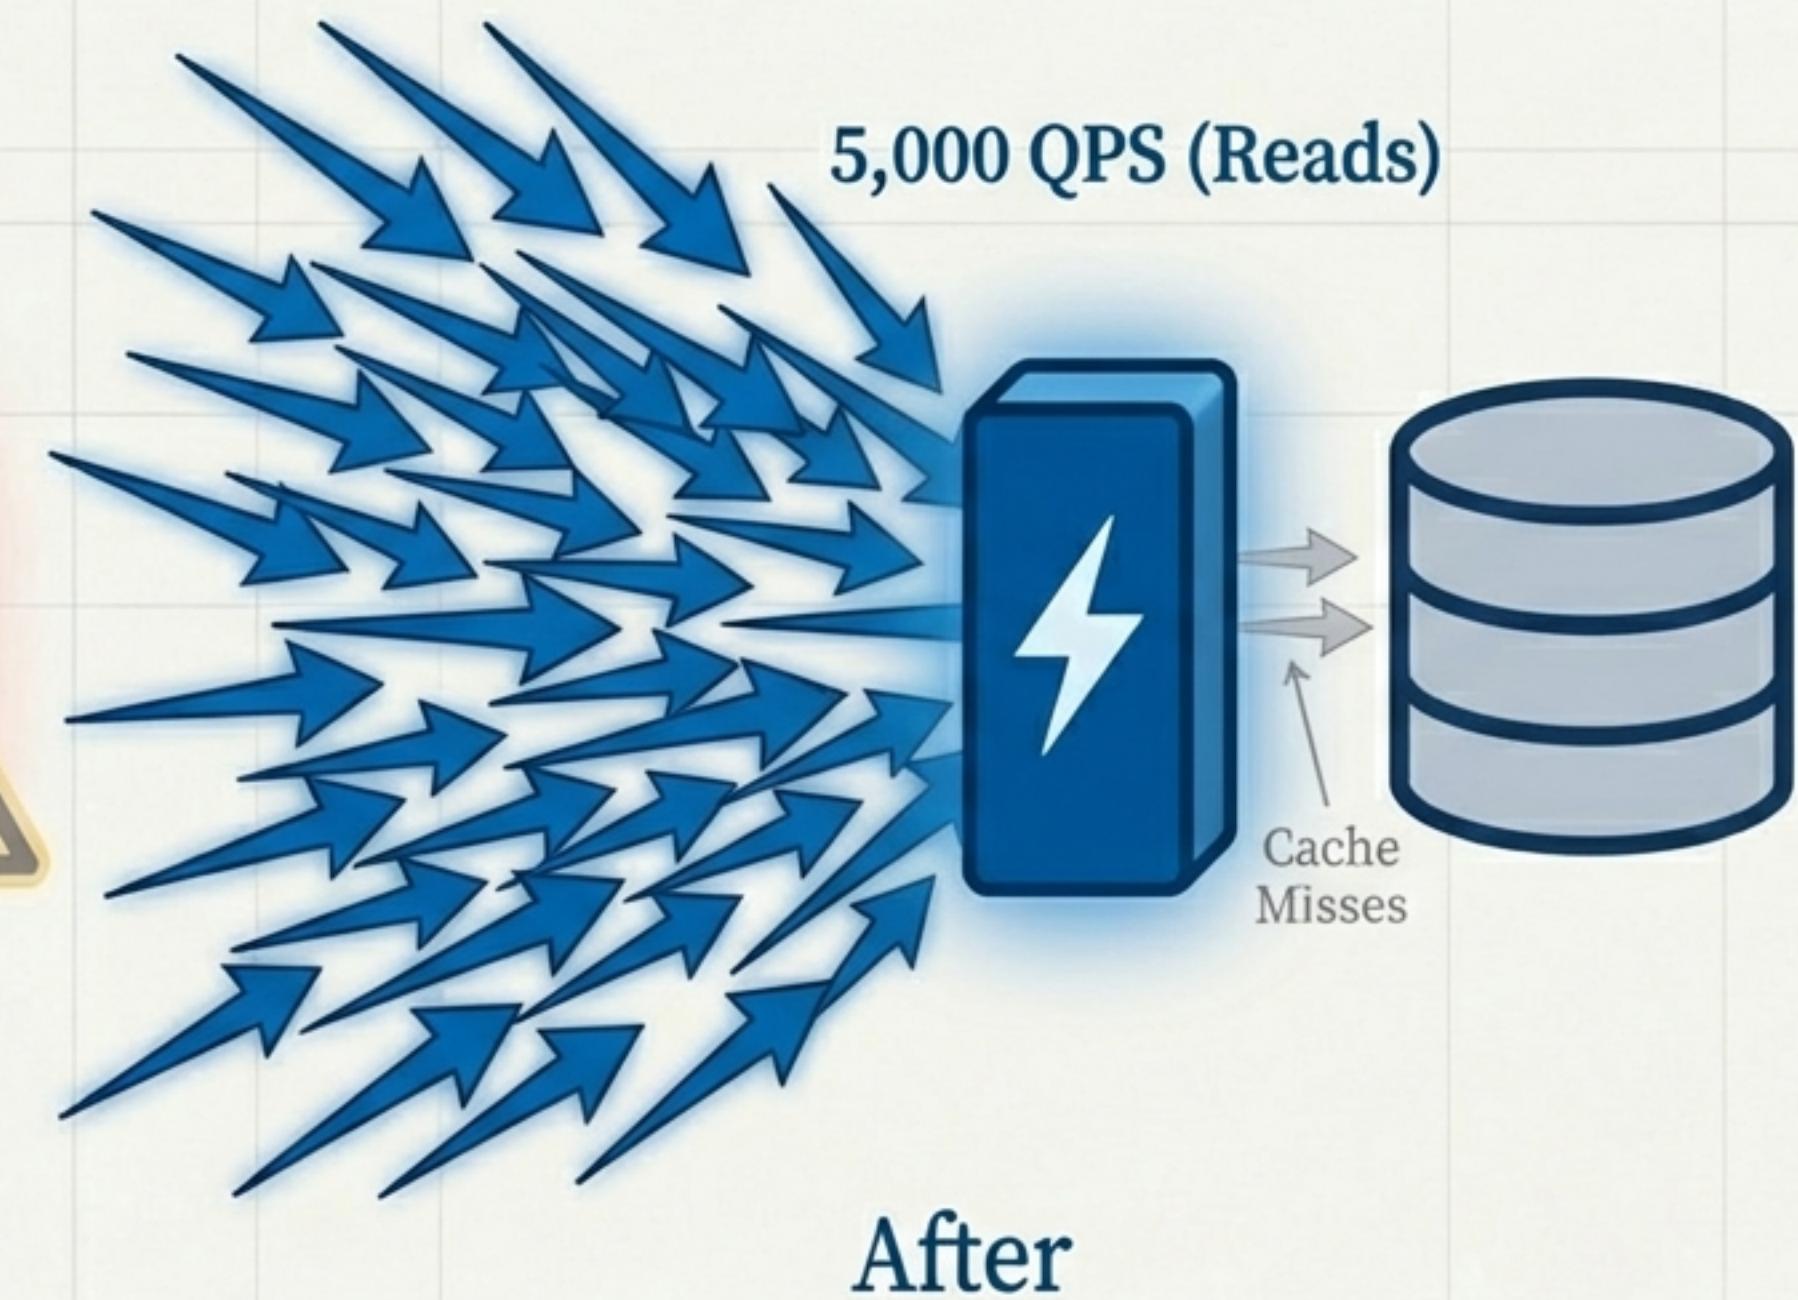
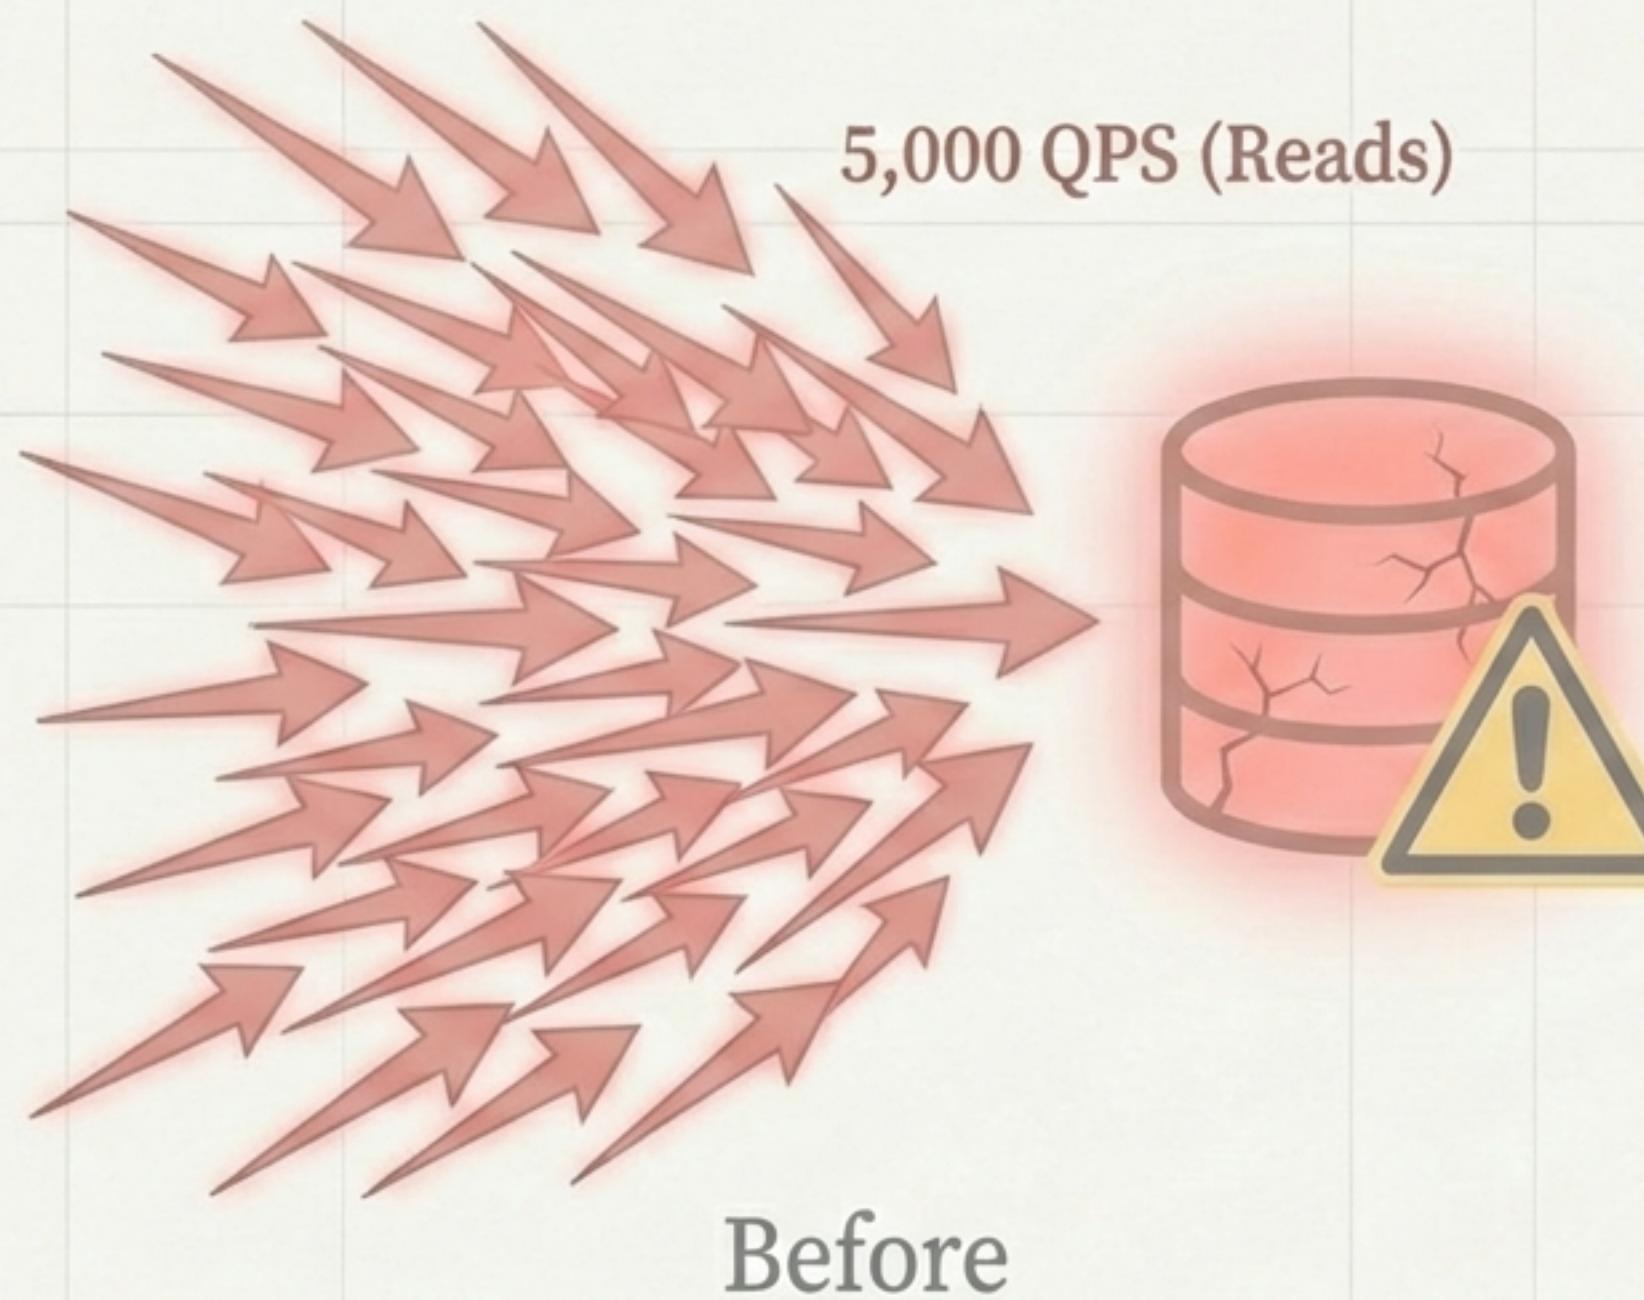


This approach leads to system failure:

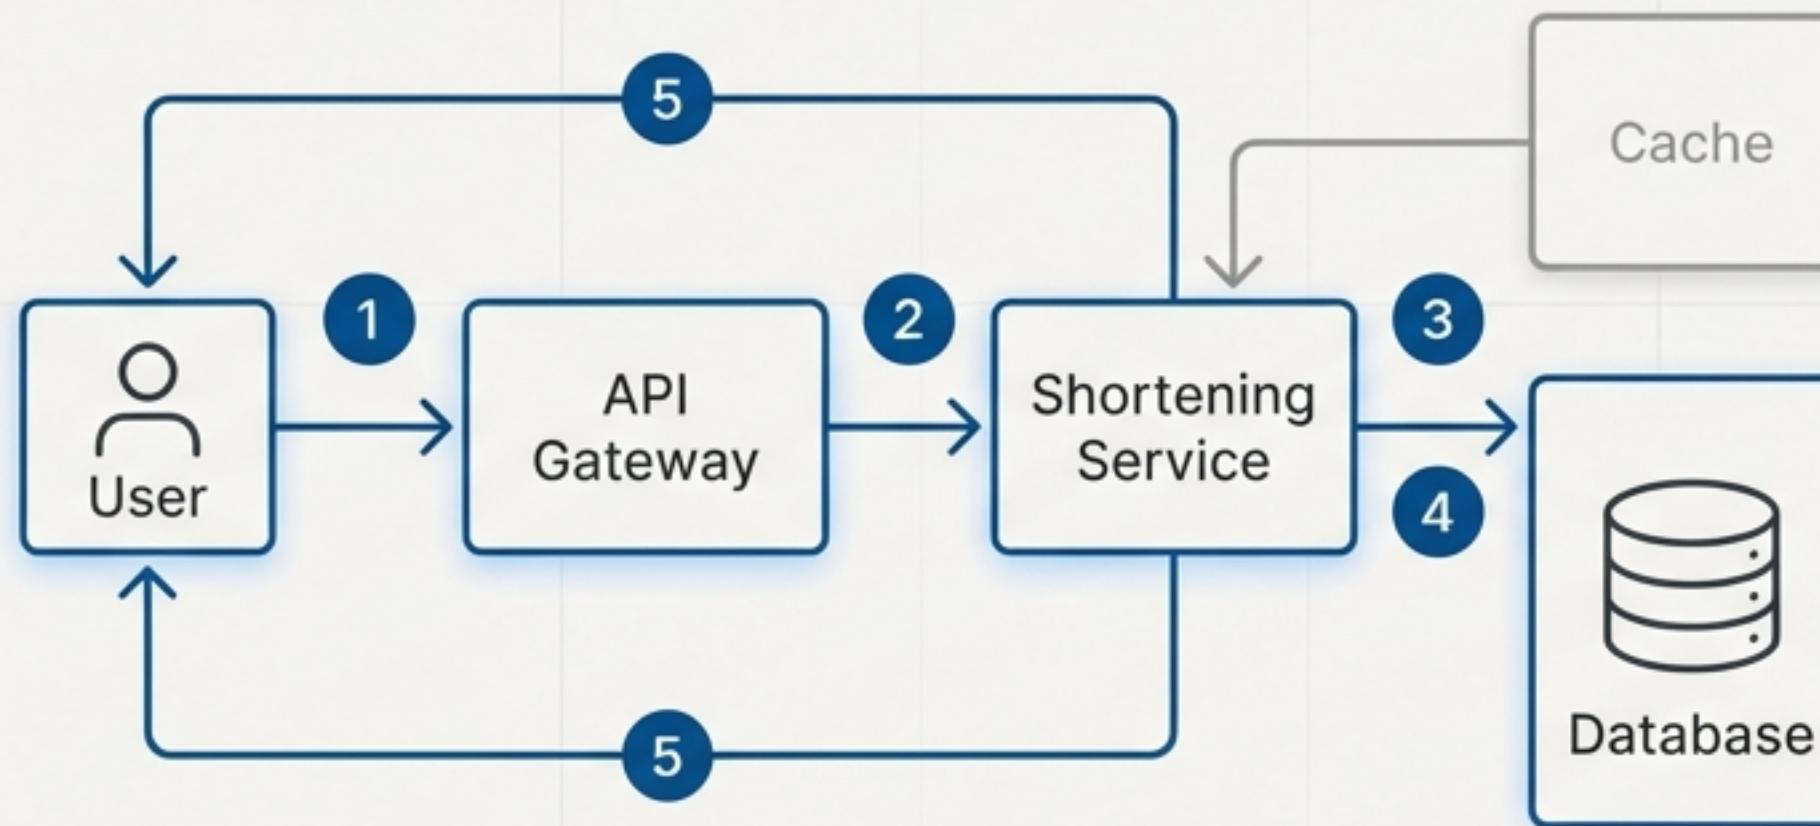
- **High Latency:** Database lookups are slow under load.
- **Exhausted Connections:** The database runs out of available connections.
- **SLA Breach:** We cannot guarantee 99.9% availability.
- **Catastrophic Failure:** The database becomes the single point of failure for the entire service.

The solution is mandatory, not optional.

**Redirects *must* be served from an in-memory cache.**

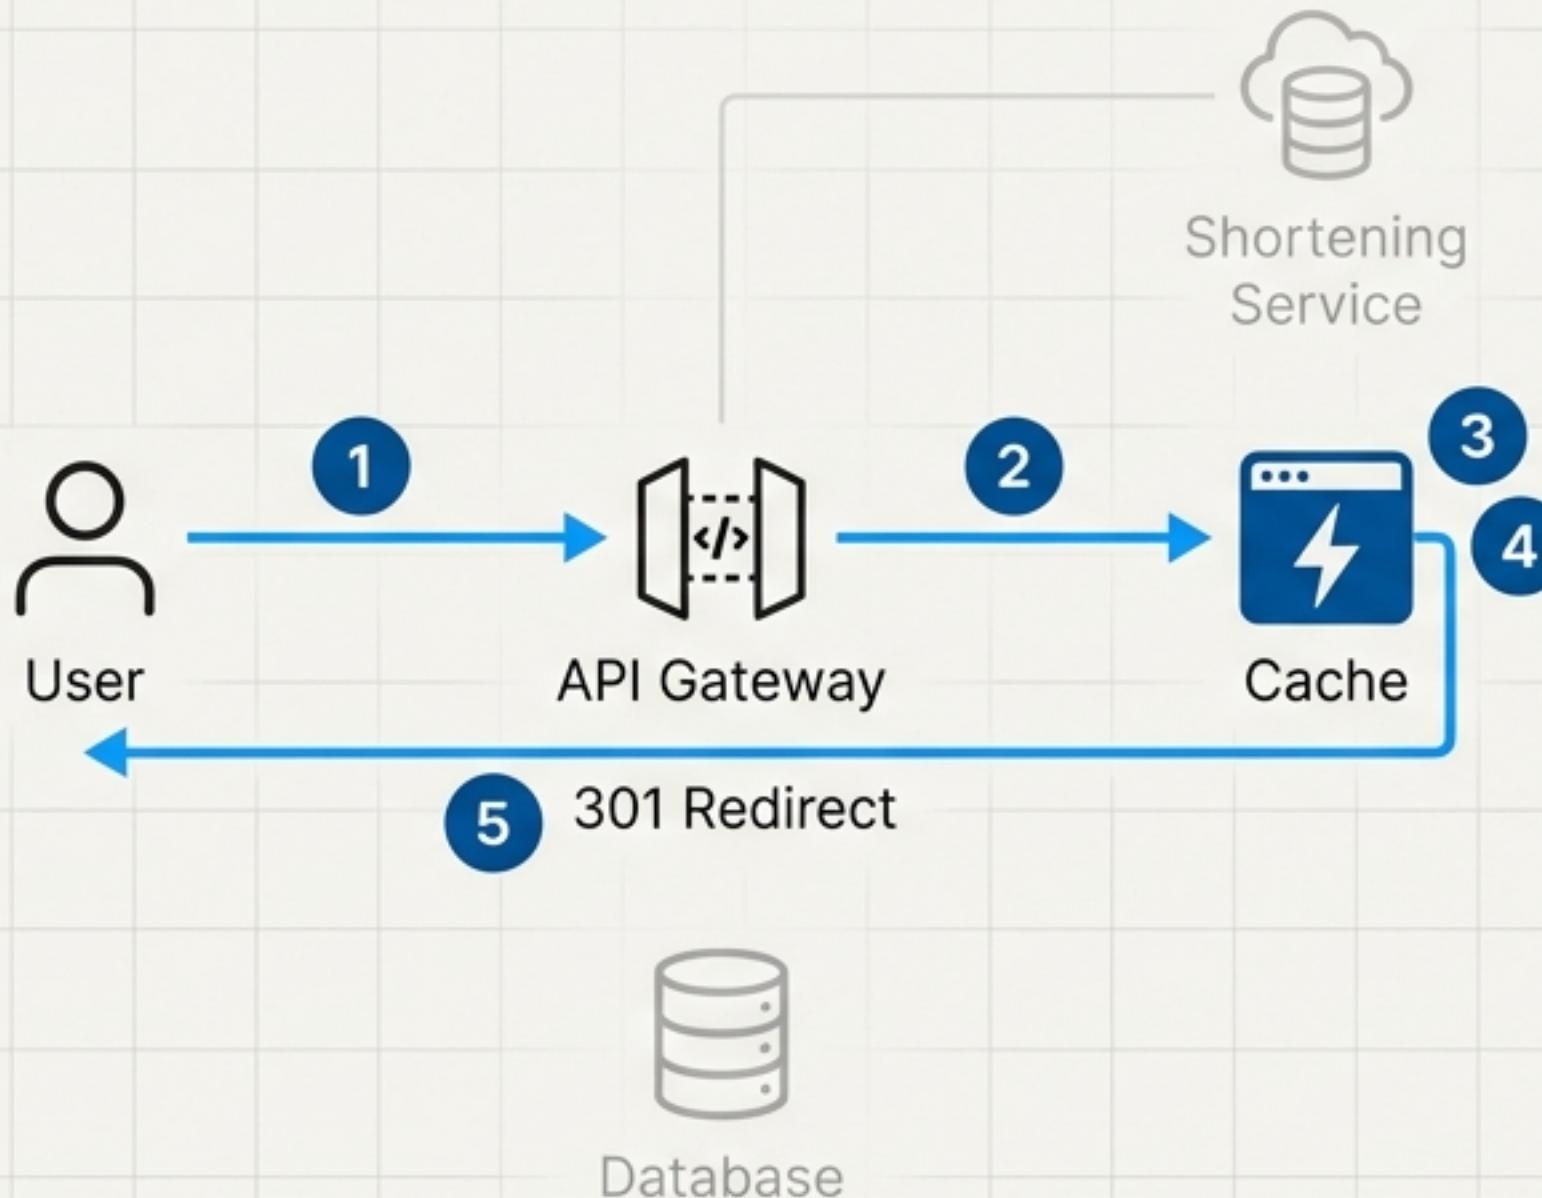


# The Final Blueprint (Part 1): The Write Path



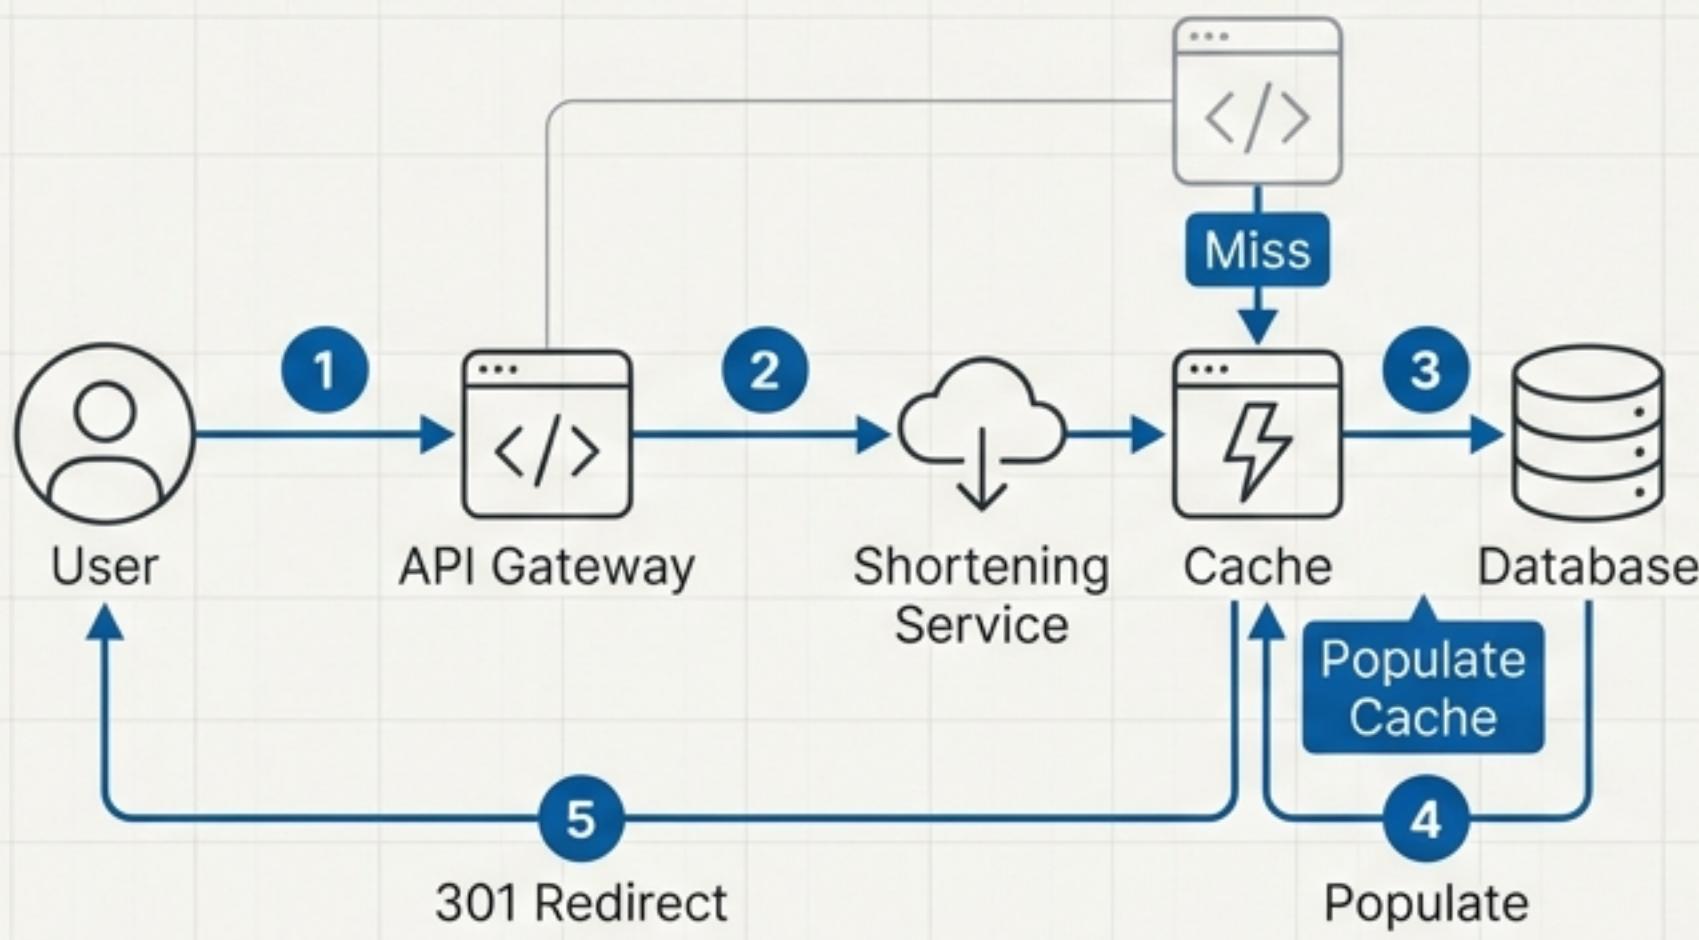
- 1 A user sends a `POST` request with the original URL.
- 2 The API Gateway routes the request to the Shortening Service.
- 3 The service generates a unique short code.
- 4 The (short\_code, original\_url) pair is written to the primary database.
- 5 The new short URL is returned to the user.

# The Final Blueprint (Part 2): The Read Path (Cache Hit)



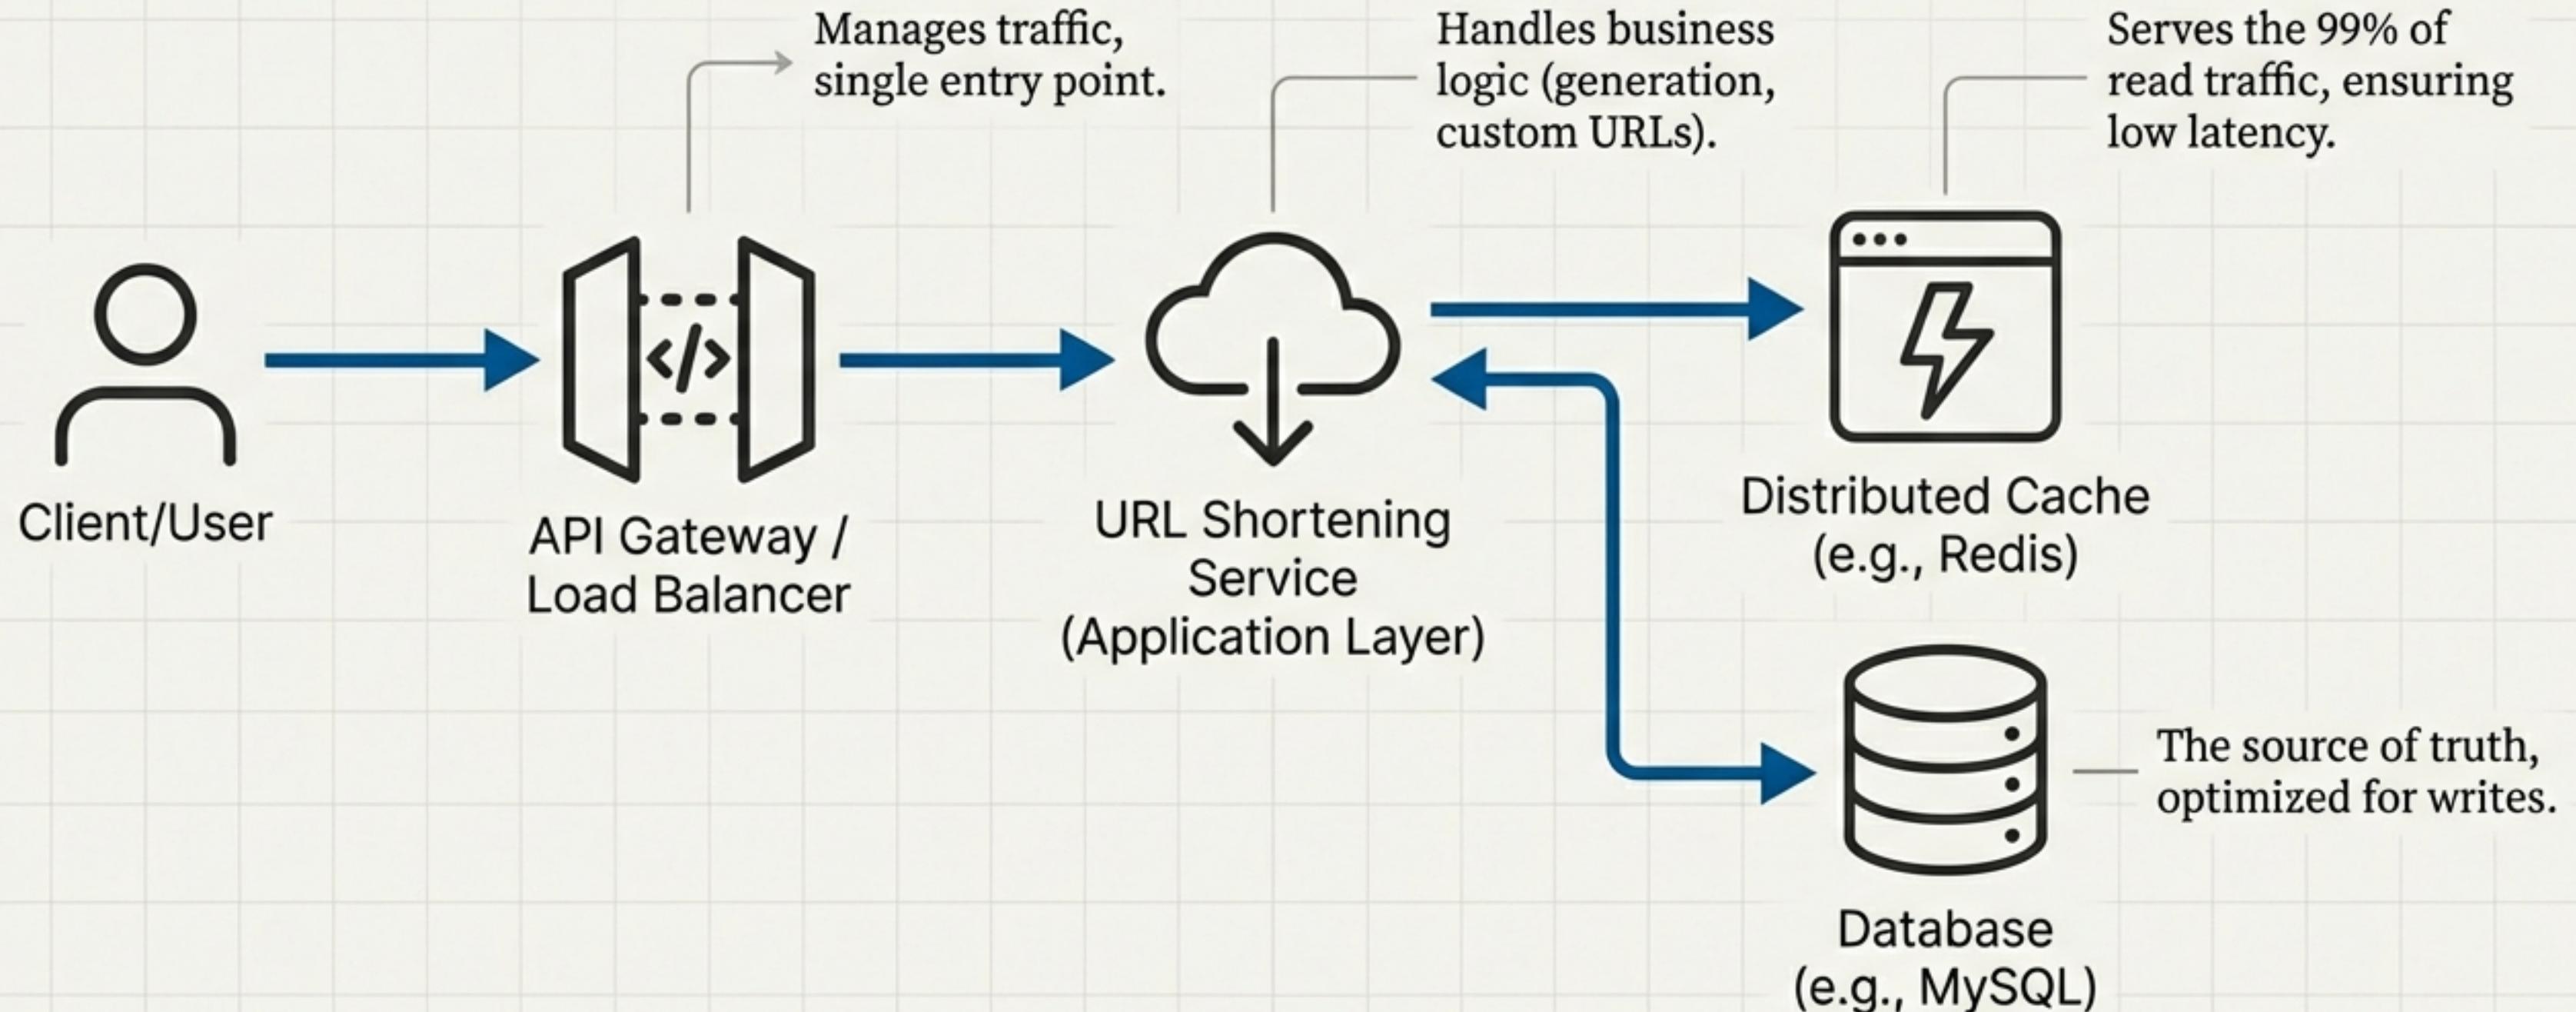
1. A user clicks a short link (**GET** request).
2. The API Gateway routes the request.
3. The system checks the Cache for the short code.
4. **Cache Hit!** The original URL is found instantly in memory.
5. A 301 redirect is returned to the user. **This is the fastest possible path.**

# The Final Blueprint (Part 3): The Read Path (Cache Miss)

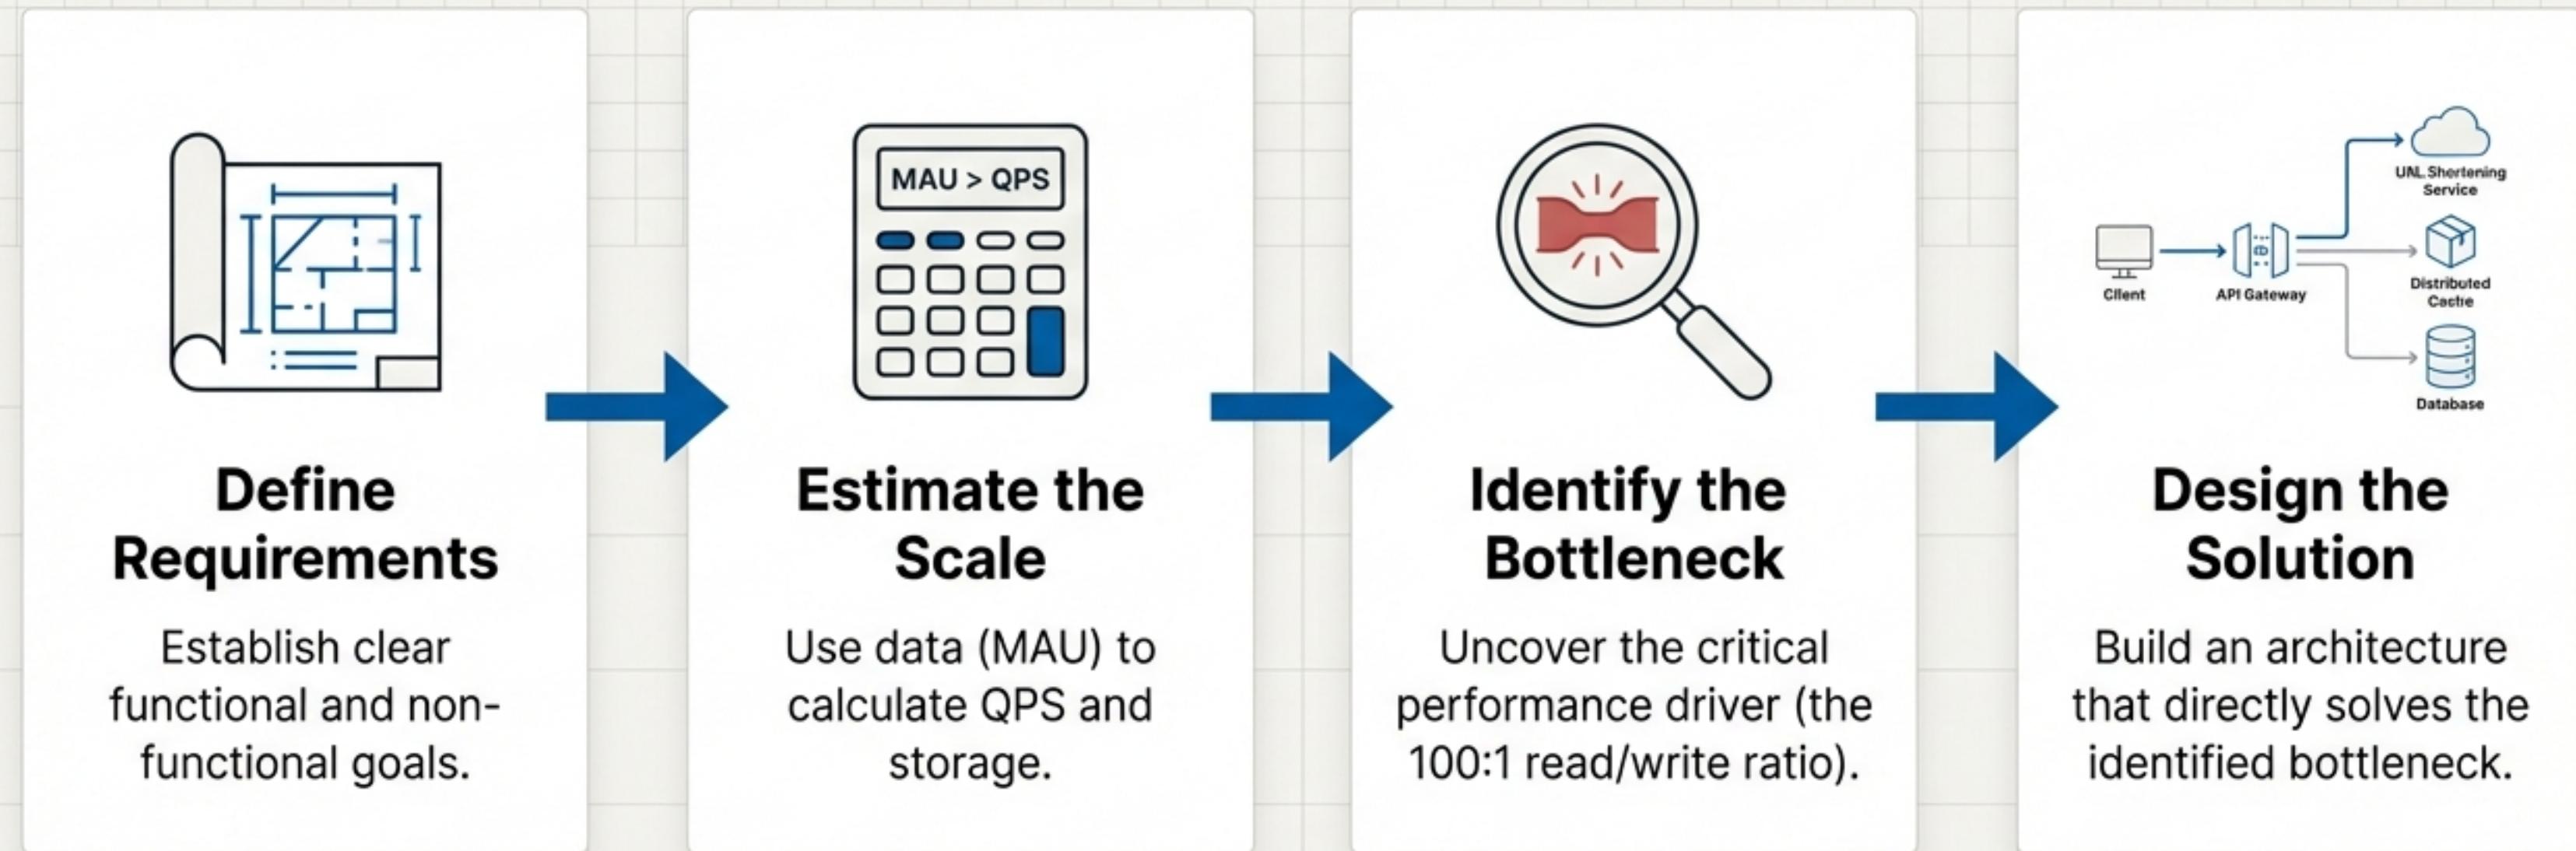


1. The system checks the Cache, but the code is not found (**Cache Miss**).
2. The system queries the primary Database for the short code.
3. The database returns the original URL.
4. **Crucial Step:** The (short\_code, original\_url) pair is written back to the Cache for future requests.
5. A 301 redirect is returned to the user.

# Our Final Architecture: Built for Scale and Speed



# The Architect's Process: From Problem to Solution



A robust system isn't just built; it's justified.